

Composable Extensions (CX) Task Group Proposal

Draft 2024-02-15

Jan Gray, Guy Lemieux
Soft CPU SIG

Outline

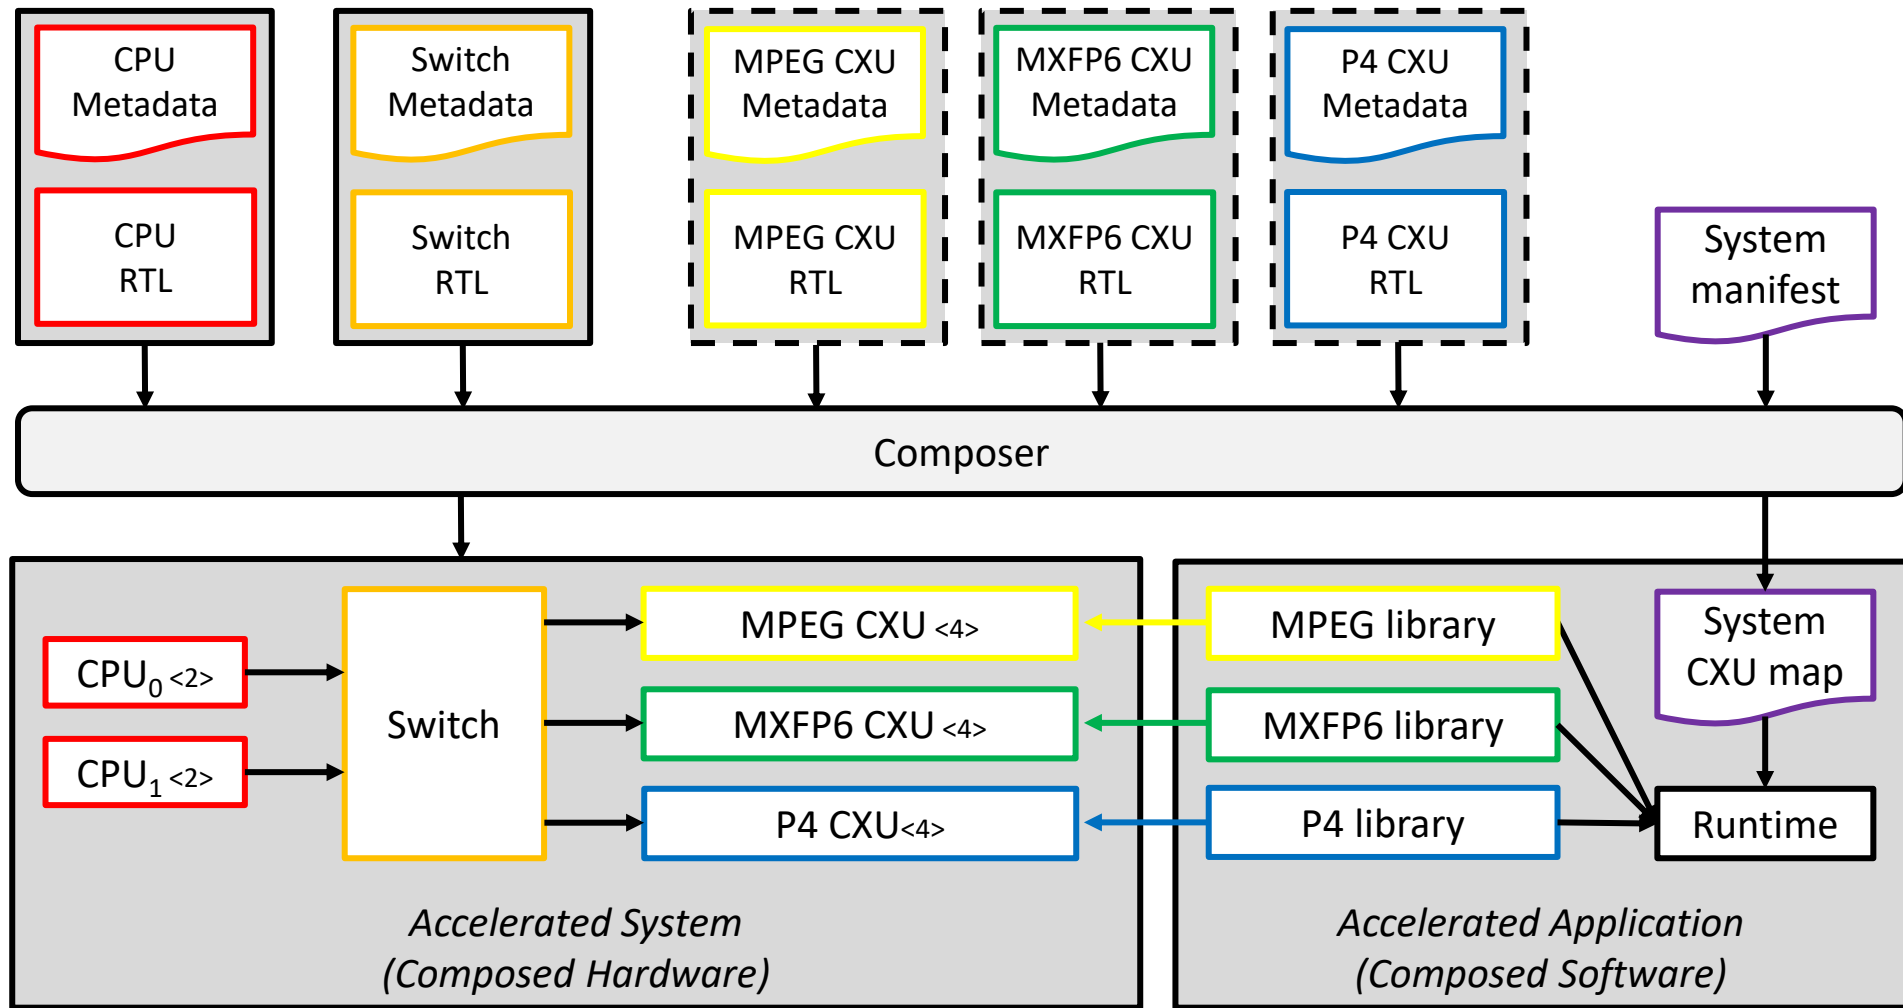
- TG ask, value proposition, support, charter
- How it might work from software
- Overview of deliverables
- Q & A

Composable Extensions TG Value Proposition

- We propose a new Task Group
- To specify ISA extensions + software and hardware interop interfaces
- To compose *my composable custom extensions* with *yours*, without prior coordination and without changing our libraries or OS
- Enables **robust reuse** of anyone's extensions and libraries
- Provides **uniform** software access to such extensions
- For a **marketplace** of extensions & libraries (especially on FPGAs)

↑ uniformity ↑ reuse ↑ agility ↓ conflicts ↓ silos ↓ fragmentation

Composing Custom Extensions



Composable Extensions TG Charter

<https://github.com/riscv-admin/sig-soft-cpu/blob/main/TG/CX/CHARTER.md>

Executive Summary

This Charter governs the TG for a Composable Extensions (CX) framework enabling **robust composition** of multiple independently authored composable custom extensions, alongside legacy custom extensions, assembled conflict-free in one RISC-V system.

By multiplexing the custom opcode space, and adopting common software and hardware interop interfaces, CX enables **uniform** extension naming, discovery, and versioning, error handling, state context management, extension hardware reuse, and stable software binaries for each target system -- all without a central management authority.

Support for the Task Group

Meta (RISC-V Summit NA 2023)

3 - Standard Specification for Customization

RISC-V considered custom extensions from its inception

- Reserved a portion of the encoding space for custom instructions

Various vendors provide different customization capabilities, but:

- Vast divergence between “what” can be customized → From nothing (most vendors), to almost everything (e.g., Codaip's Coda)
- No unified way to specify custom extensions → Vendor specific tools, interfaces and capabilities
- No common way to introduce these extensions to software

Provide a standard way to specify custom extensions:

- Specify architecturally visible features, such as custom registers and instructions
- Unify how these extensions are introduced to software
- Leave room for vendors to innovate at the microarchitecture and implementation level

CX TG supporters/participants

- AMD (Xilinx)
- Lattice
- Bluespec
- Antmicro
- Alibaba
- Individual members
- *More to come*

Composable Extension Multiplexing

For conflict-free reuse of the custom opcode space

- Multiplexing off: legacy custom instructions, unchanged
- Extension library:
 - **Discover**: is named extension available here? → CX selector *index*
 - **Select**: `csrrw prev-index, cx_index, index`
 - **Issue**: `custom-*` instructions
- CPU:
 - **Decode**: `custom-*` instruction → (function-ID, operands)
 - **Issue**: send request (state-ID, function-ID, [insn,] operands) to *selected* extension
 - **Retire**: receive response, update dest-register and `cx_status` CSR

Stateful Composable Extensions

- CX custom instructions may be stateful
 - May access accumulators, registers, register files, ...
 - Some CX state accessible as *CX scoped CSRs (CX-CSRs)*
- Uniform CX state context management
 - Supports 100s of CX state contexts, per CX, per system
 - Any mapping from harts to CX state contexts, 1:1, 1: n , n :1
 - The same OS code can save/restore any state context of any CX

Example CX Programming Model, in C++

```
struct MyEncoderCX : CX {
    static cx_guid_t guid;          // unique canonical ID of MyEncoderCX contract
    // CX for stateless sum of absolute differences, 8x8b SIMD inputs, 64b output
    static uint64_t sad_cx(uint64_t a, uint64_t b) { return cxfn(1, a, b); }
    static uint64_t sad_sw(uint64_t a, uint64_t b); // software version
    ...
};

// CX library example: accelerated sum of absolute differences of two vectors
static uint64_t MyEncoderCX::sad_vecs(uint64_t as[], uint64_t bs[], size_t n) {
    CX::Scope encoder(guid);        // discover & select CX if present; writes cx_index
    uint64_t total = 0;
    for (size_t i = 0; i < n; ++i) {
        // use CX instruction if present, software otherwise;
        // here 'encoder' test is loop invariant, can be hoisted
        total += encoder ? sad_cx(as[i], bs[i]) : sad_sw(as[i], bs[i]);
    }
    return total;
    // encoder.~Scope() automatically restores caller's CX selection; writes cx_index
}
```

Design Tenets

- **Routine robust reuse of composable extensions & libraries**
- No prior coordination or central authority
- Stable software binaries
- Stable modular hardware
- Uniform SW: naming, discovery, versioning, state, errors, ...
- Simple, frugal, fast
- Secure
- Interop across the decades

CX TG Extensions and Interface Spec Deliverables

- New ISA extension(s): **new custom CSRs & multiplexed custom instructions** only

	Unpriv and Priv ISA
CX Multiplexing	mcx_selector cx_status (basic muxing) + mcx_table cx_index (optional privileged access control)
CX State Management	<i>CX scoped</i> custom CSRs + + 3 CX CSRs for OS to read/write any CX state context

- New interfaces
 - Software: CX-API: extension services: uniform discovery, versioning, state mgmt
 - Software: CX-ABI: software rules for CX-Mux CSRs
 - Hardware: CXU-LI (logic interface): auto-composition of CPUs + CX Units

Composable Extensions Collateral

TG Charter:

<https://github.com/riscv-admin/sig-soft-cpu/blob/main/TG/CX/CHARTER.md>

Draft Spec:

<https://raw.githubusercontent.com/grayresearch/CX/main/spec/spec.pdf>

Design and Rationale Talk:

https://www.youtube.com/watch?v=7daY_E2itpo

Talk slides:

<https://raw.githubusercontent.com/grayresearch/CFU/main/collateral/design-rationale-CX-CXU-spec.pdf>

Repo: <https://github.com/grayresearch/CX>

Draft Proposed RISC-V
Composable Custom Extensions
Specification

Version 0.92.231111, 2023-11-11: Draft

Task Group Personnel

- Past contributors
 - Tim Ansell, Tim Callahan, Jan Gray, Karol Gugala, Olof Kingdren, Maciej Kurc, Guy Lemieux, Charles Papon, Tim Vogt
(Antmicro, Lattice, Google, Gray Research, QAMCOM, SpinalHDL/VexRiscv, UBC)
- TG interest
 - AMD, Lattice, Bluespec, AntMicro, Alibaba, individuals
- TG Chair/Vice-Chair Interest
 - Jan Gray, Guy Lemieux, TBD / open invite

Conclusion / Q & A