

# Composable Custom Extensions and Custom Function Units for RISC-V

Jan Gray (Gray Research) , Tim Vogt (Lattice Semiconductor), Tim Callahan (Google), Charles Papon (SpinalHDL), Guy Lemieux (University of British Columbia), Maciej Kurc (Antmicro), Karol Gugala (Antmicro) **draft #2, 2022-04-26**



## RISC-V extensions allow diverse uses. But fragmentation!

- Standard extensions layer and compose, by design
  - But consume finite encoding space, take years to ratify
- Custom extensions enable rapid in-house accelerator+library solutions
  - FPGA RISC-V SoCs: *insight* → *accelerated SoC* in an hour
  - But custom extensions may not work *together* in a system, due to conflicting encodings, CPU pipelines, logic interfaces, models of discovery, computation, state, error handling, versioning, tooling
  - Siloed solutions hurt HW & SW reuse and fragment the ecosystem

## A mix-and-match custom extensions ecosystem

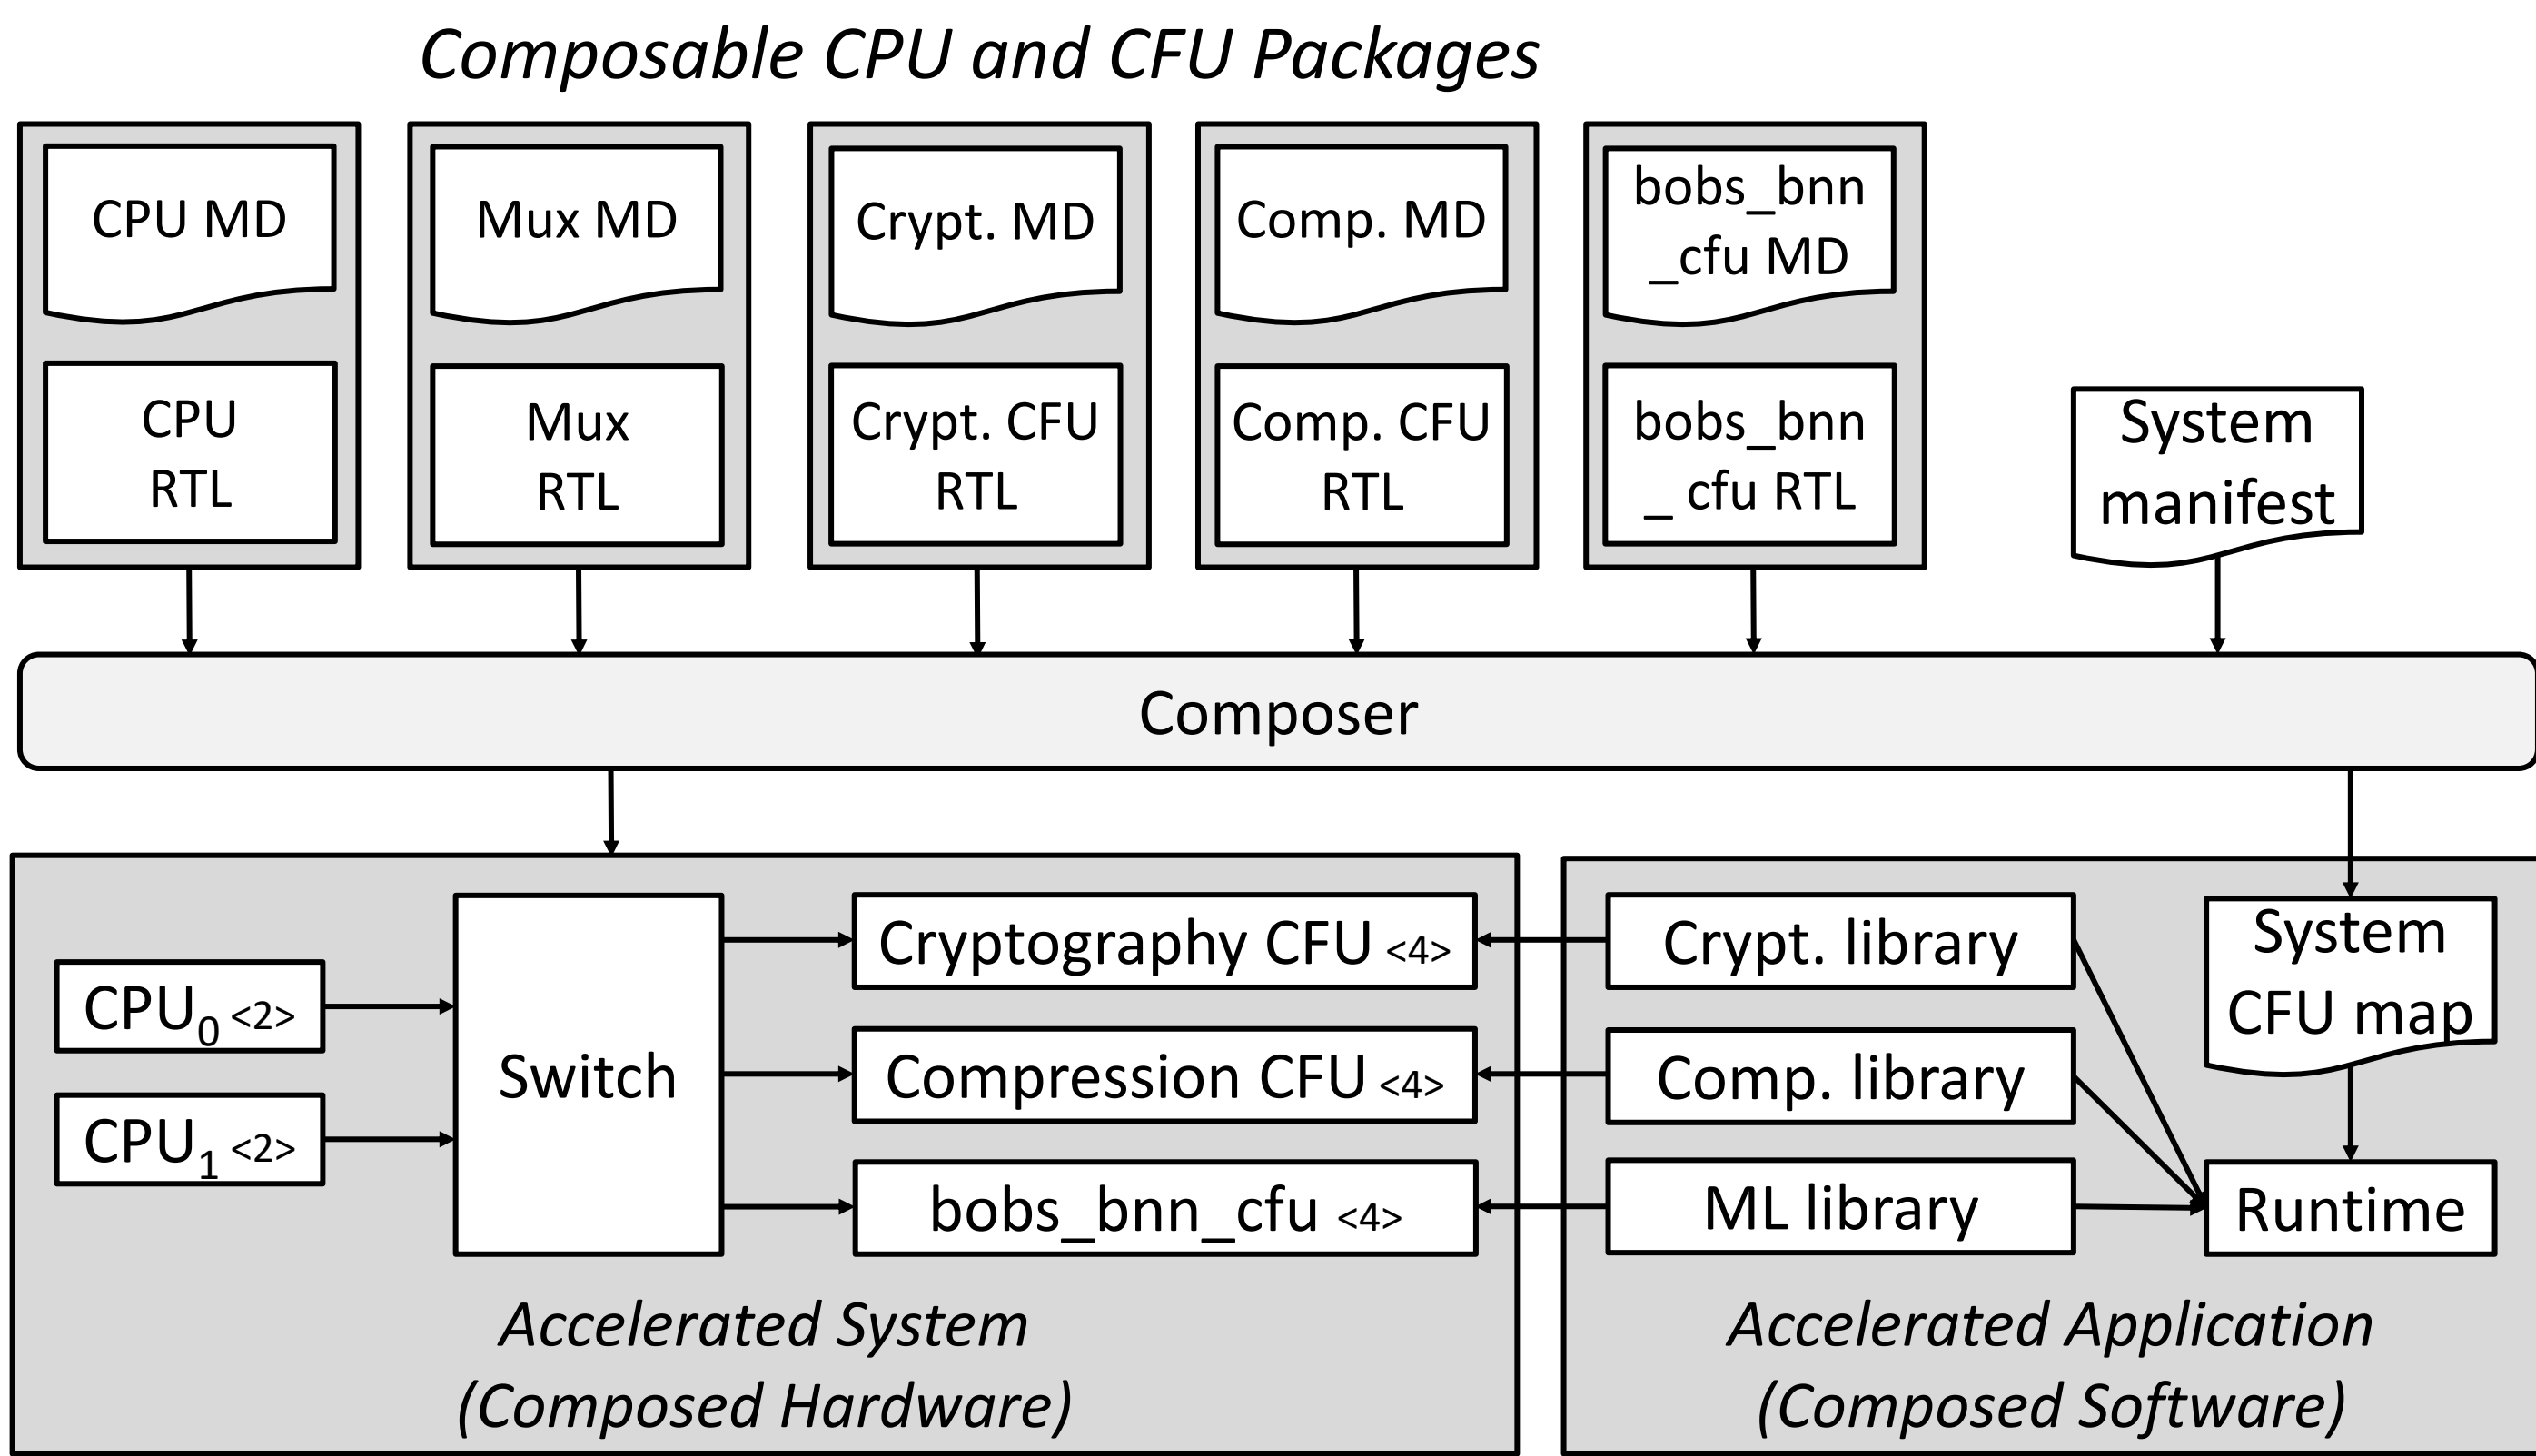
- *Agility* of custom extensions with *composability* of standard extensions
- Proposed hardware and software interop interfaces enable a new ecosystem of reusable accelerators & libraries that *just work* together

## Custom interface, custom function unit, accelerated library

- *Custom interface (CI)*: a composable custom extension: a named, immutable set of composition-safe custom instructions
- *Custom function unit (CFU)*: a core that implements a custom interface
- *Accelerated library*: issues custom instructions to a custom interface
- Different CFUs may implement an interface; different libs may use one
- Anybody can create a custom interface, CFU, or library, a CPU that composes CFUs, and tools for same, such that all works with all, always

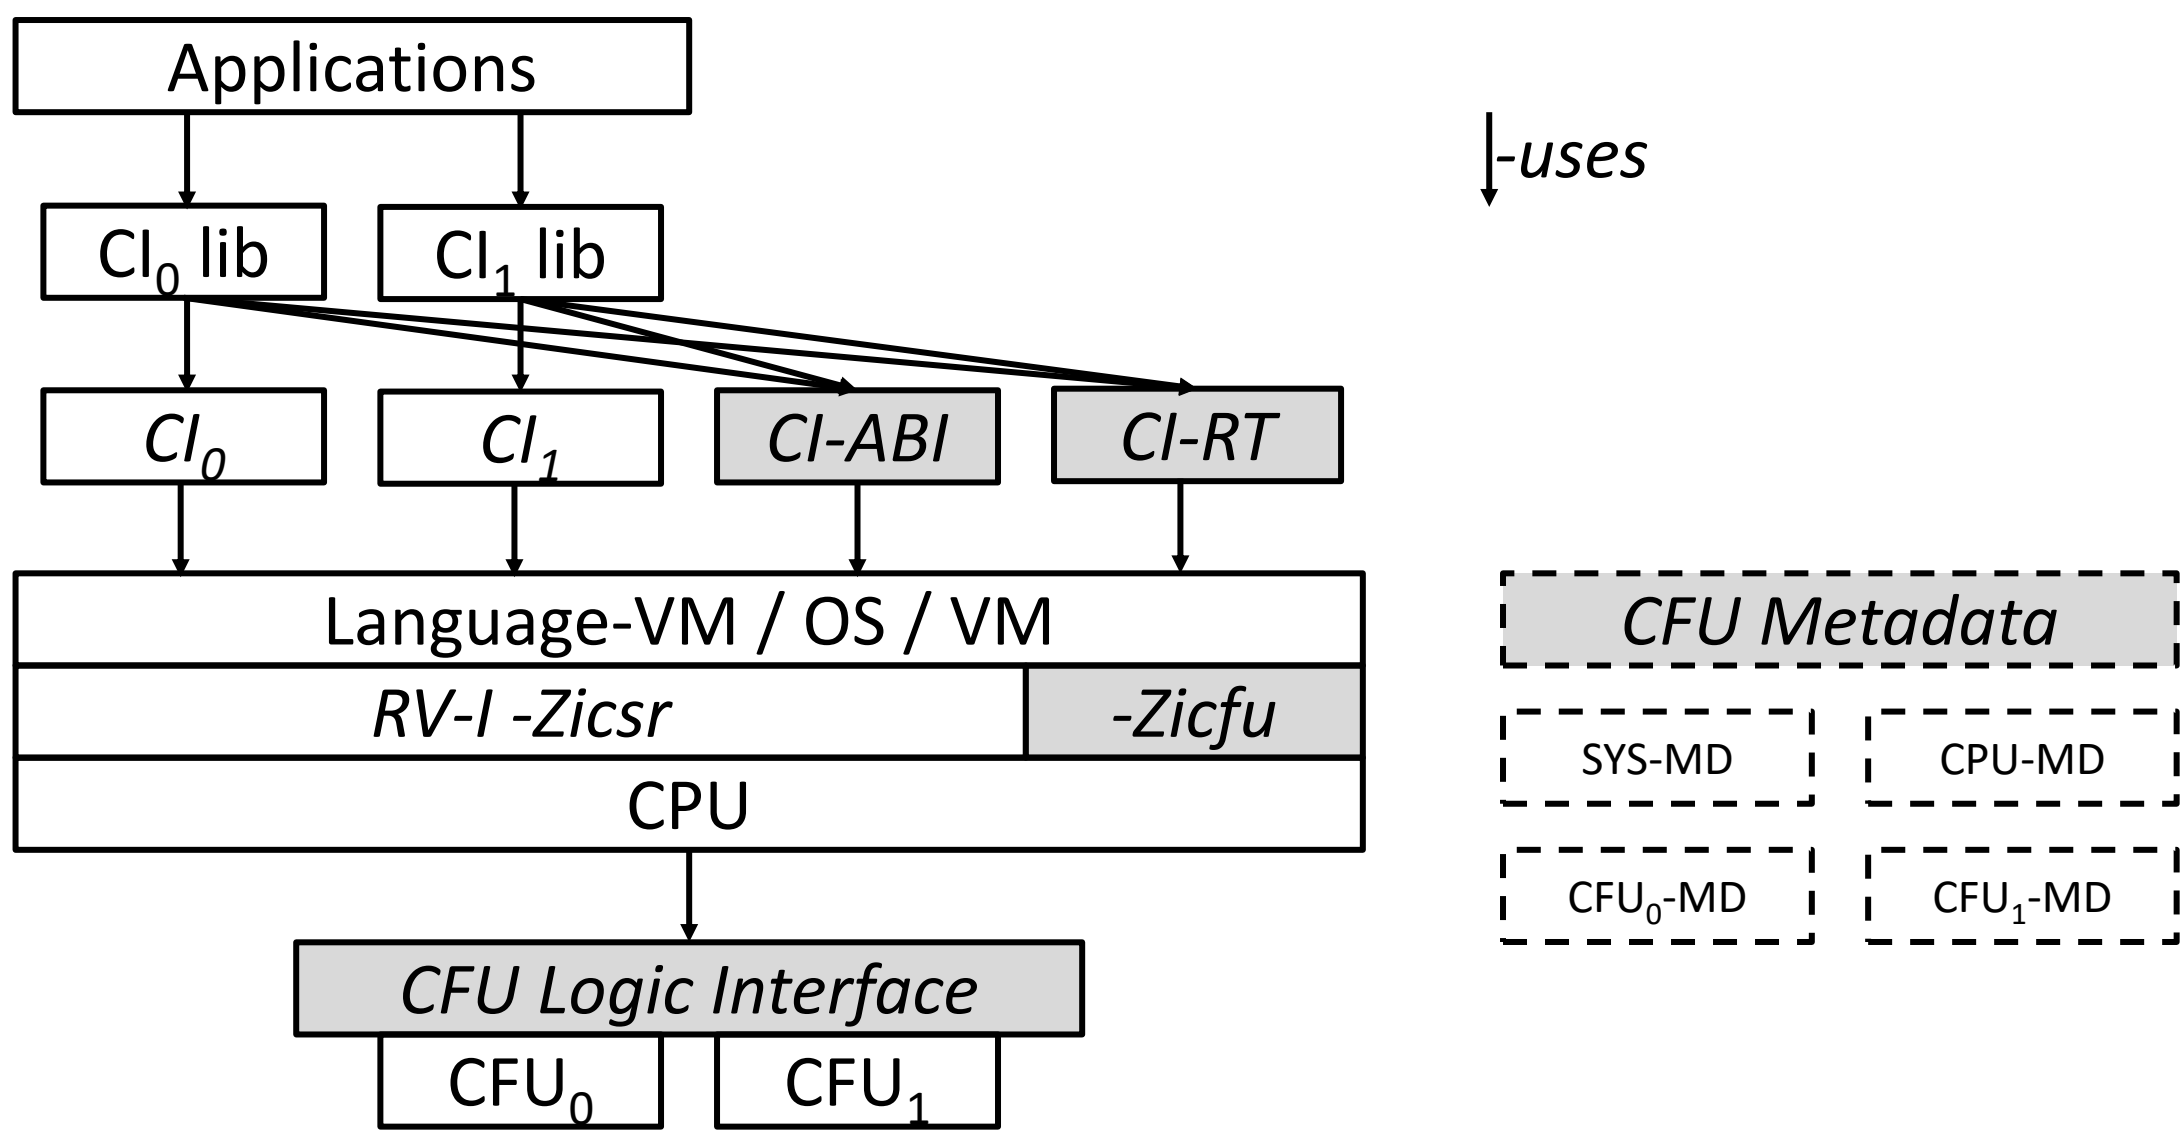
## Example

- Alice’s MPSoC uses a custom cryptography library and its CFU
- The compression lib can use custom instructions too. Alice adds its CFU
- Bob adds his new binary neural net custom interface, ML lib, and CFU
- Using Runtime, each lib *selects* its interface, then issues its instructions



## HW-SW stack incorporating proposed interop interfaces

- Custom Interface ABI, Custom Interface Runtime
- “-Zicfu”: CSRs: `mcfu_select`, `cfu_status`; instructions: `custom-0/1/2`
- CFU Logic Interface: for automatic CPUs + CFUs composition



## Composition challenges → our proposed solutions

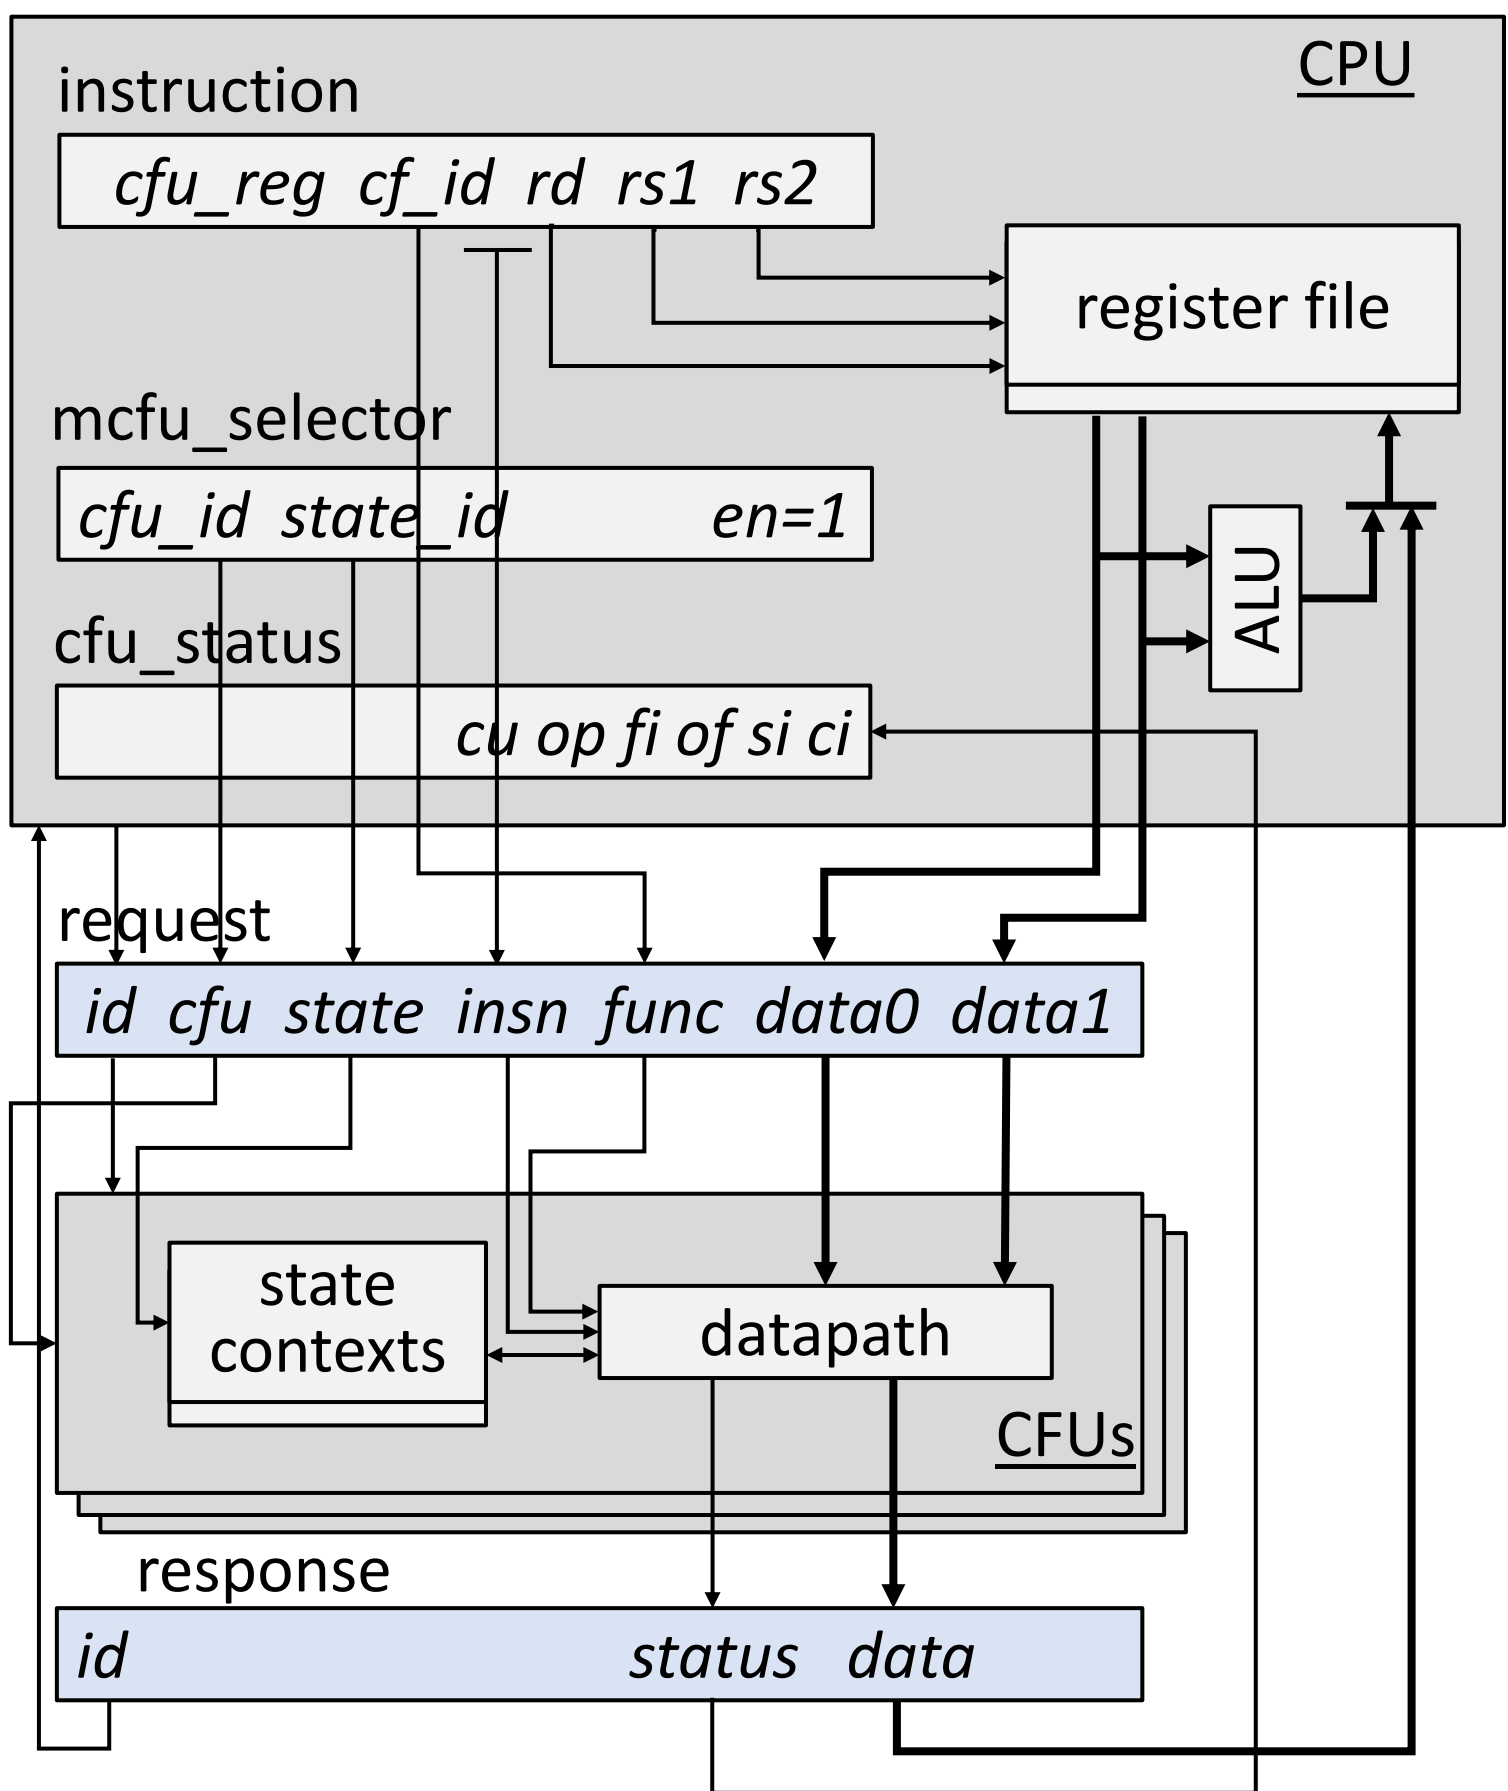
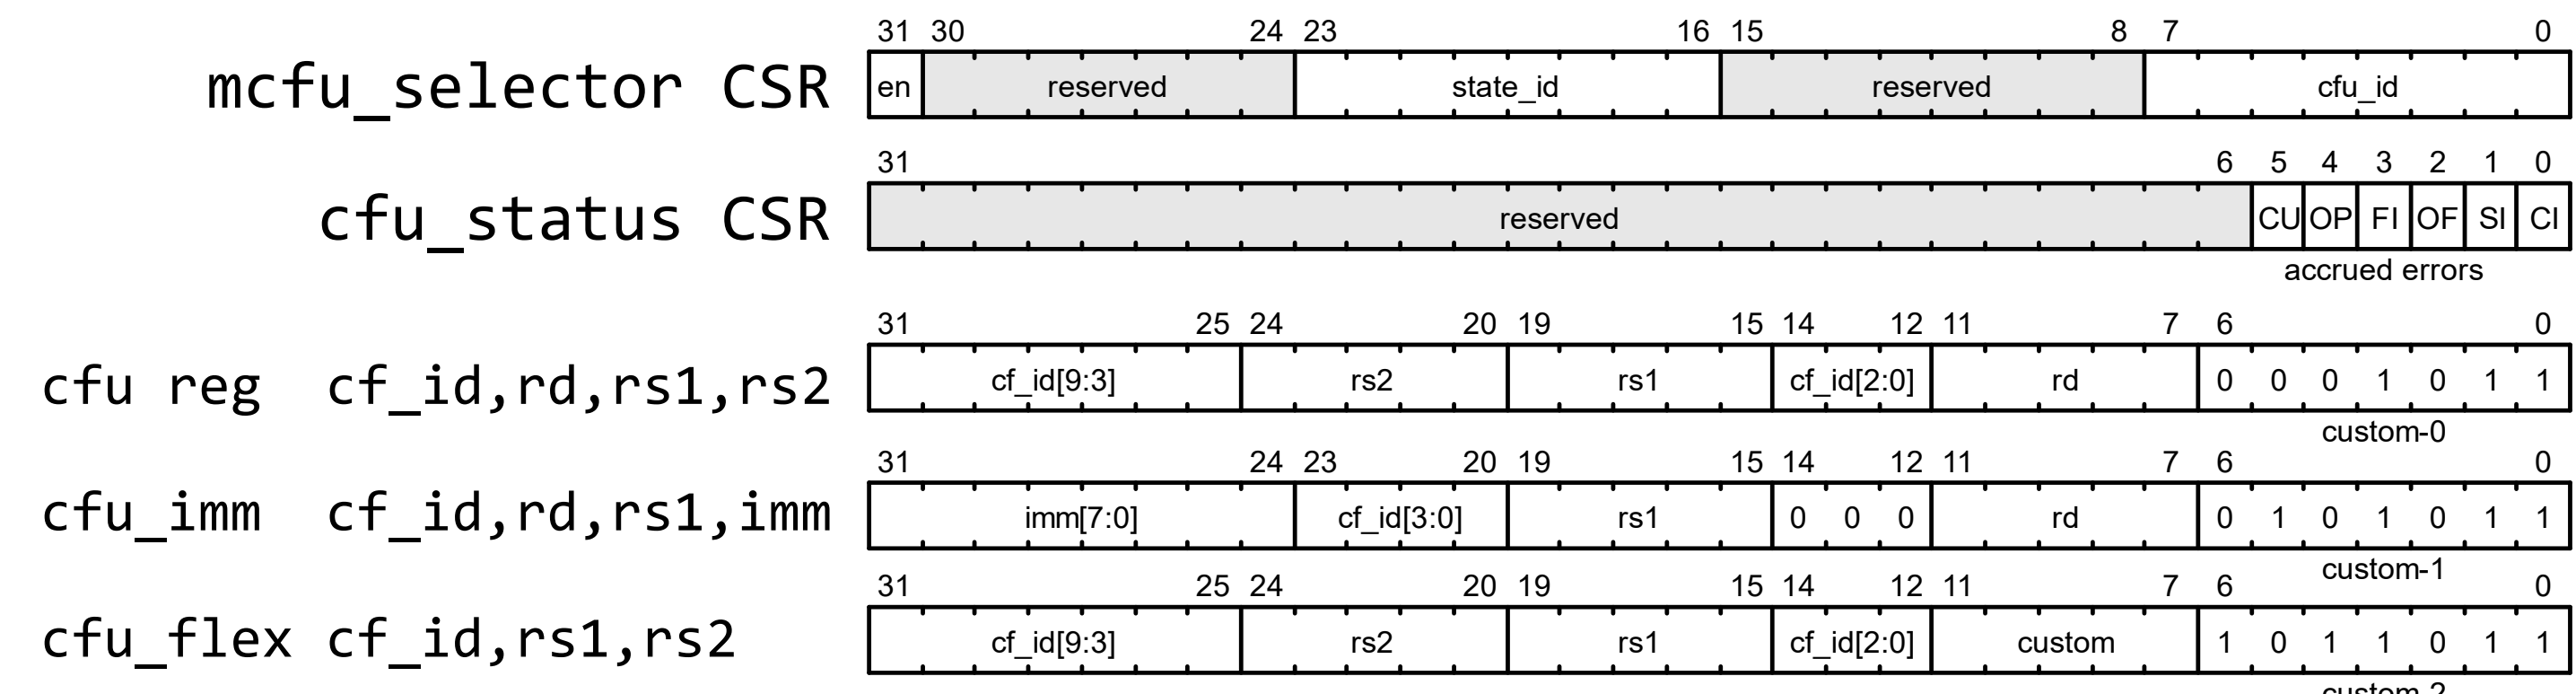
- Naming? Everything scoped to unique 128b-GUID custom interface IDs
- Opcode space? Inexhaustible with custom interface multiplexing
- Correctness under composition? Strict isolation of interfaces’ states
- Error handling? `mcfu_status` CSR accrues errors
- Discovery? The Runtime consults a generated System CFU Map
- Context save/restore? Use standard *IStateContext* custom instructions
- Versioning? Immutable interfaces + newer-interface discovery
- Access control? Optional privileged access control of CFUs & state
- Please see spec for all the details

## Correct composition via strict isolation of interface state

- Behavior of interfaces **must not** change when composed with others
- Custom instructions only access register file & the selected state context
- A CFU may be configured with 0, 1, #harts, or *n* isolated state contexts
- *V1*: no access to other shared resources (memory, CSRs, PC, exceptions)

## HW-SW interface: custom interface multiplexing

- Inexhaustible, collision-free opcode space & no central opcode authority
- `mcfu_selector` CSR selects hart’s *current* CFU and its state context
- `custom-0/-1/-2` instructions send a request to the current CFU
- CFU executes request, may read & write its state context
- CFU response updates destination register and `cfu_status` CSR

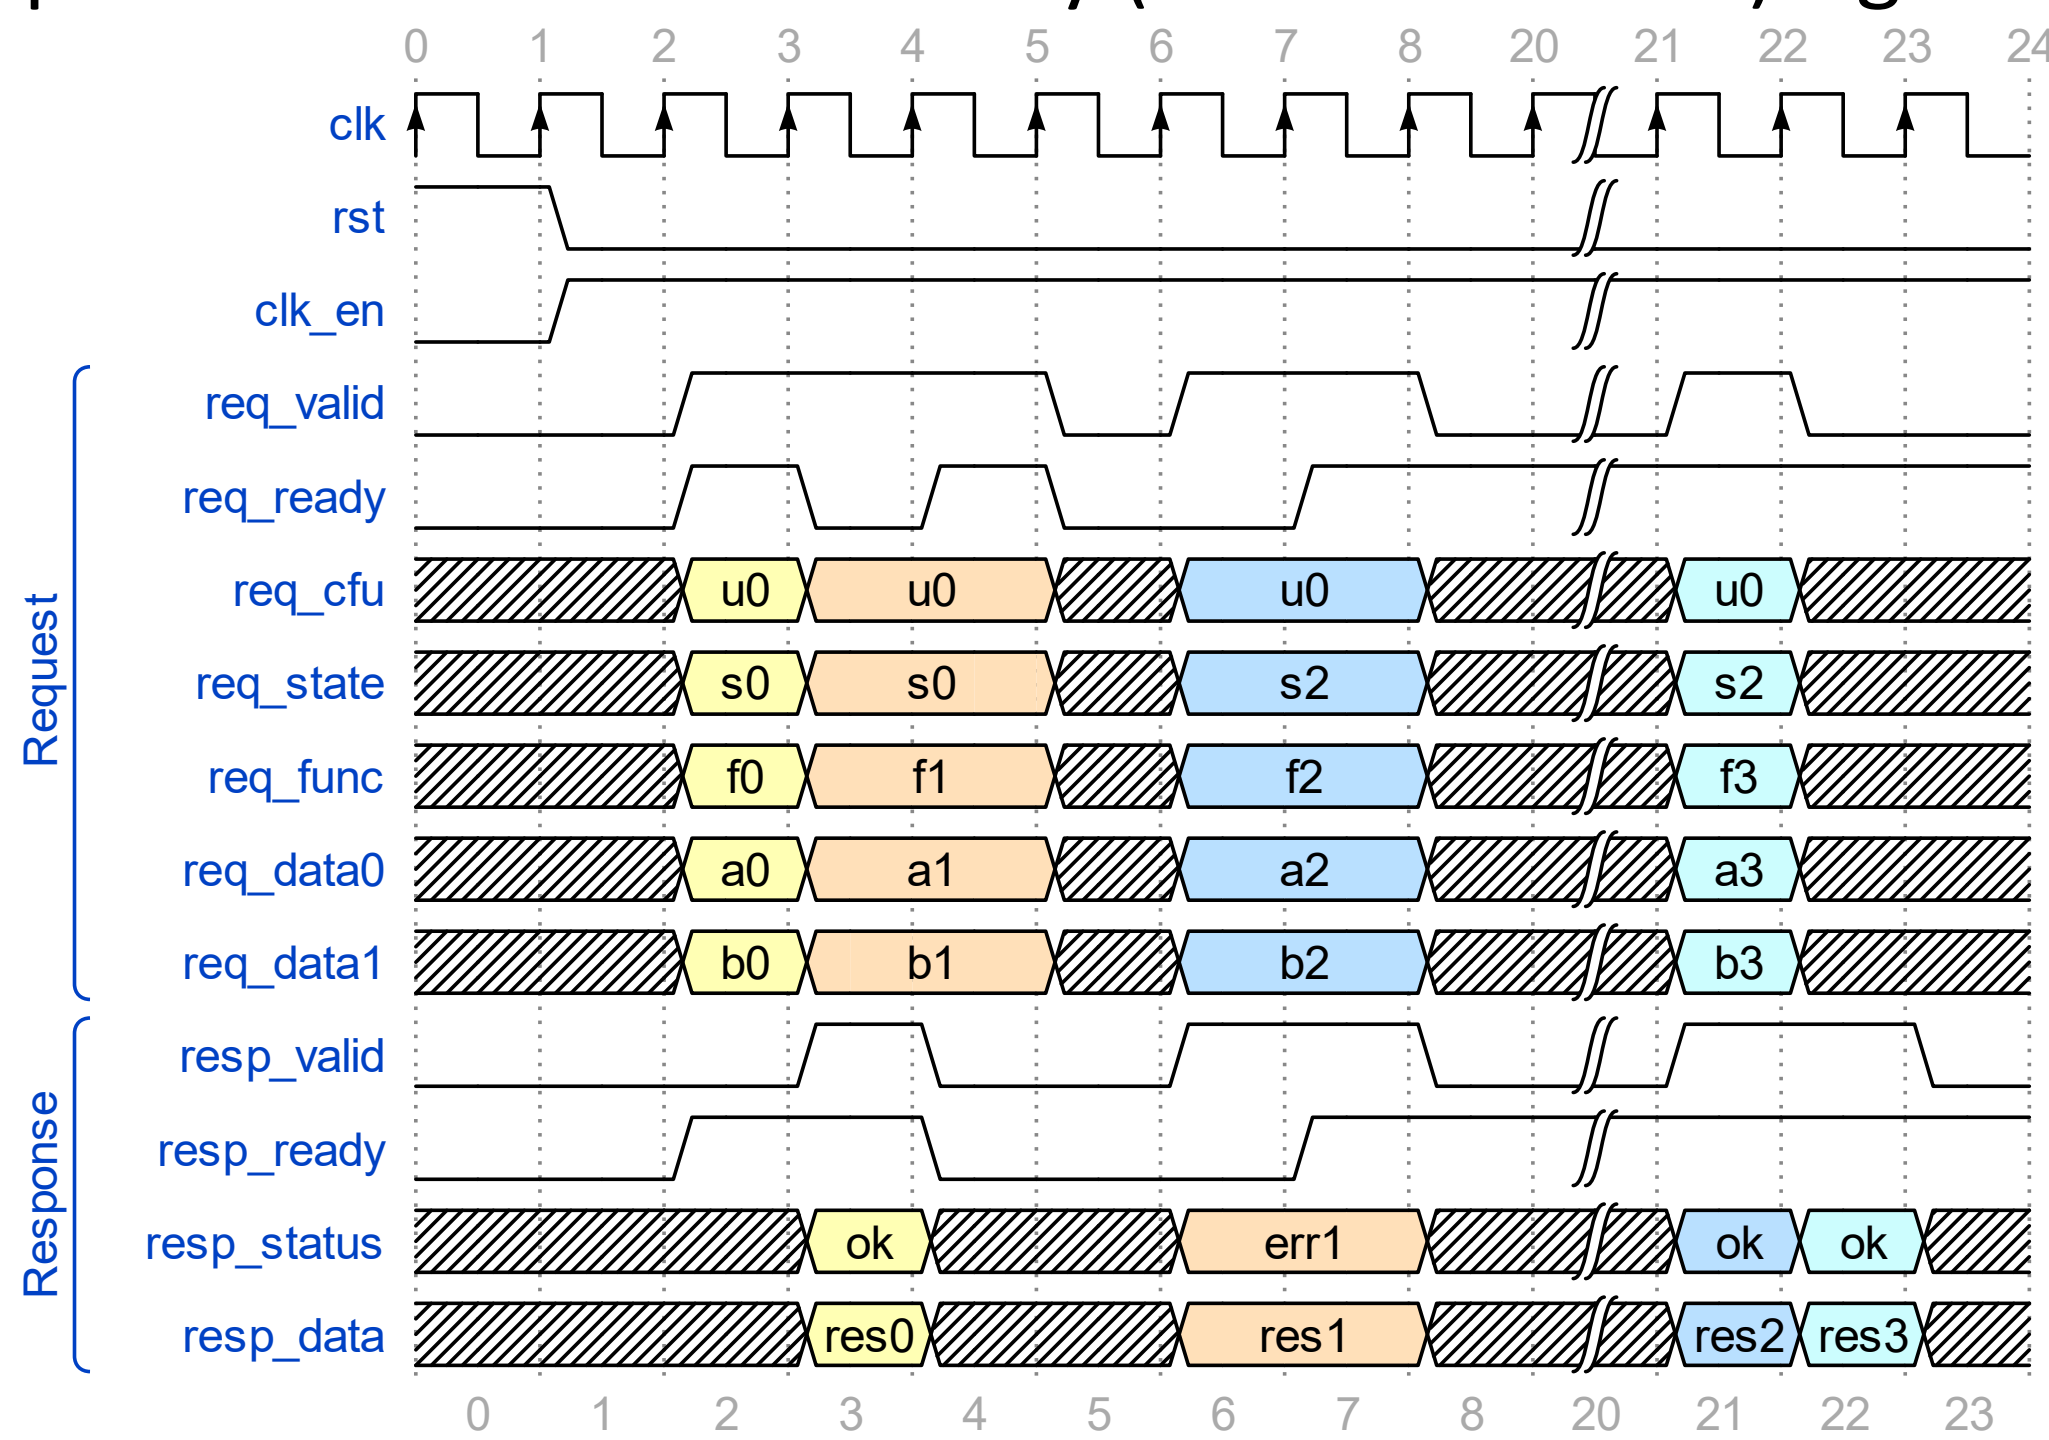


## Exemplary accelerated library programming model

- Try to select a custom interface, issue custom instructions if CFU present
- ```
CI bm(CI_ID_IBitmanip); // csrrw mcfu_selector ...
if (bm) // accelerator present:
    count = cf(pcnt_cf, data, 0); // cfu_reg pcnt_cf,rd,data,x0
else // accelerator not present:
    count = popcount(data); // software version
```

## HW-HW interface: CFU Logic Interface (CFU-LI)

- Interoperable CPUs-to-CFUs-to-CFUs signaling and metadata
- Flexible feature levels: combinational, fixed, variable latency, reordering
- Includes prebuilt switches & adapters for automatic composition
- Example: CFU-L2 variable latency (flow-controlled) signaling:



## Status, more information, help us shape these interfaces

- **Draft Proposed RISC-V Composable Custom Extensions Specification, <https://github.com/grayresearch/CFU>**
- Status: refining spec, building consensus, writing RTL to demo mix-and-match composition across various CPUs and multiple CFUs
- Discussions and meetings: see RISC-V [sig-soft-cpu] list

**Let us add just enough architecture  
that our custom extensions just work together**